



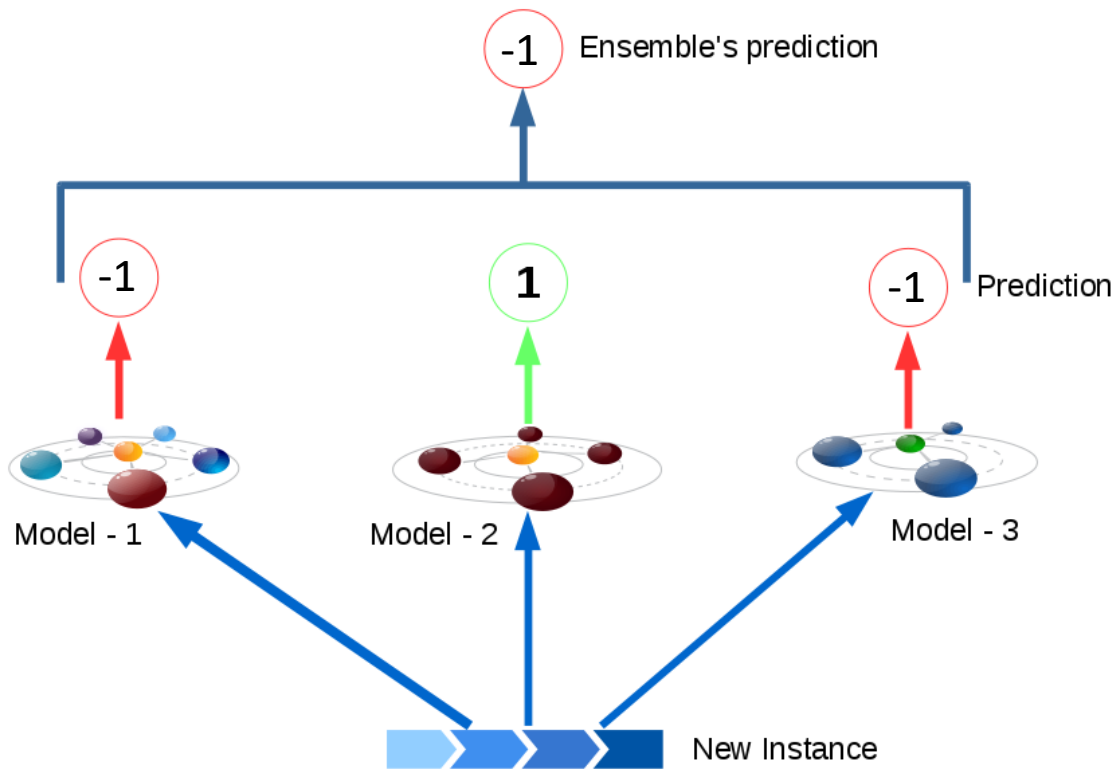
Introduction to Machine Learning

**Ensemble Methods,
Bagging, and Boosting**

Power through diversity



Multiple classifiers: Voting




Power through diversity



Let's try this out

Multiple training sets: Resampling

- Sampling with replacement



\tilde{D}

Multiple training sets: Bagging

- Start with training data set D with N examples
- Use sampling with replacement to create M datasets $\tilde{D}_1, \dots, \tilde{D}_M$
 - Each has size N
 - Train separate classifier on each training set
 - Combine (vote)

Boosting weak learners

- Weak learners become strong through adaptation
- AdaBoost




Algorithm 32 $\text{ADABOOST}(\mathcal{W}, \mathcal{D}, K)$

```
1:  $d^{(0)} \leftarrow \langle \frac{1}{N}, \frac{1}{N}, \dots, \frac{1}{N} \rangle$  // Initialize uniform importance to each example
2: for  $k = 1 \dots K$  do
3:    $f^{(k)} \leftarrow \mathcal{W}(\mathcal{D}, d^{(k-1)})$  // Train  $k$ th classifier on weighted data
4:    $\hat{y}_n \leftarrow f^{(k)}(x_n), \forall n$  // Make predictions on training data
5:    $\hat{\epsilon}^{(k)} \leftarrow \sum_n d_n^{(k-1)} [y_n \neq \hat{y}_n]$  // Compute weighted training error
6:    $\alpha^{(k)} \leftarrow \frac{1}{2} \log \left( \frac{1 - \hat{\epsilon}^{(k)}}{\hat{\epsilon}^{(k)}} \right)$  // Compute “adaptive” parameter
7:    $d_n^{(k)} \leftarrow \frac{1}{Z} d_n^{(k-1)} \exp[-\alpha^{(k)} y_n \hat{y}_n], \forall n$  // Re-weight examples and normalize
8: end for
9: return  $f(\hat{x}) = \text{sgn} [\sum_k \alpha^{(k)} f^{(k)}(\hat{x})]$  // Return (weighted) voted classifier
```

Let's try this out

Decision Stump? Really?



Random ensembles

- Decision trees are expensive to learn
- If structure was given would be faster


Algorithm 33 RANDOMFORESTTRAIN(\mathcal{D} , $depth$, K)

```
1: for  $k = 1 \dots K$  do  
2:    $t^{(k)} \leftarrow$  complete binary tree of depth  $depth$  with random feature splits  
3:    $f^{(k)} \leftarrow$  the function computed by  $t^{(k)}$ , with leaves filled in by  $\mathcal{D}$   
4: end for  
5: return  $f(\hat{x}) = \text{sgn} [\sum_k f^{(k)}(\hat{x})]$  // Return voted classifier
```

Performance

- Early proponents of **random forests**: “Joint Induction of Shape Features and Tree Classifiers”, Amit, Geman and Wilder, PAMI 1997


Features: arrangement of **tags**



Arrangements: 8 angles

tags

Common 4x4 patterns



A subset of all the 62 tags

#Features: $62 \times 62 \times 8 = 30,752$

Single tree: **7.0%** error

Random forest of 25 trees: **0.8%** error

Let's try this out