Introduction to Machine Learning

Learning to Rank

Ranking

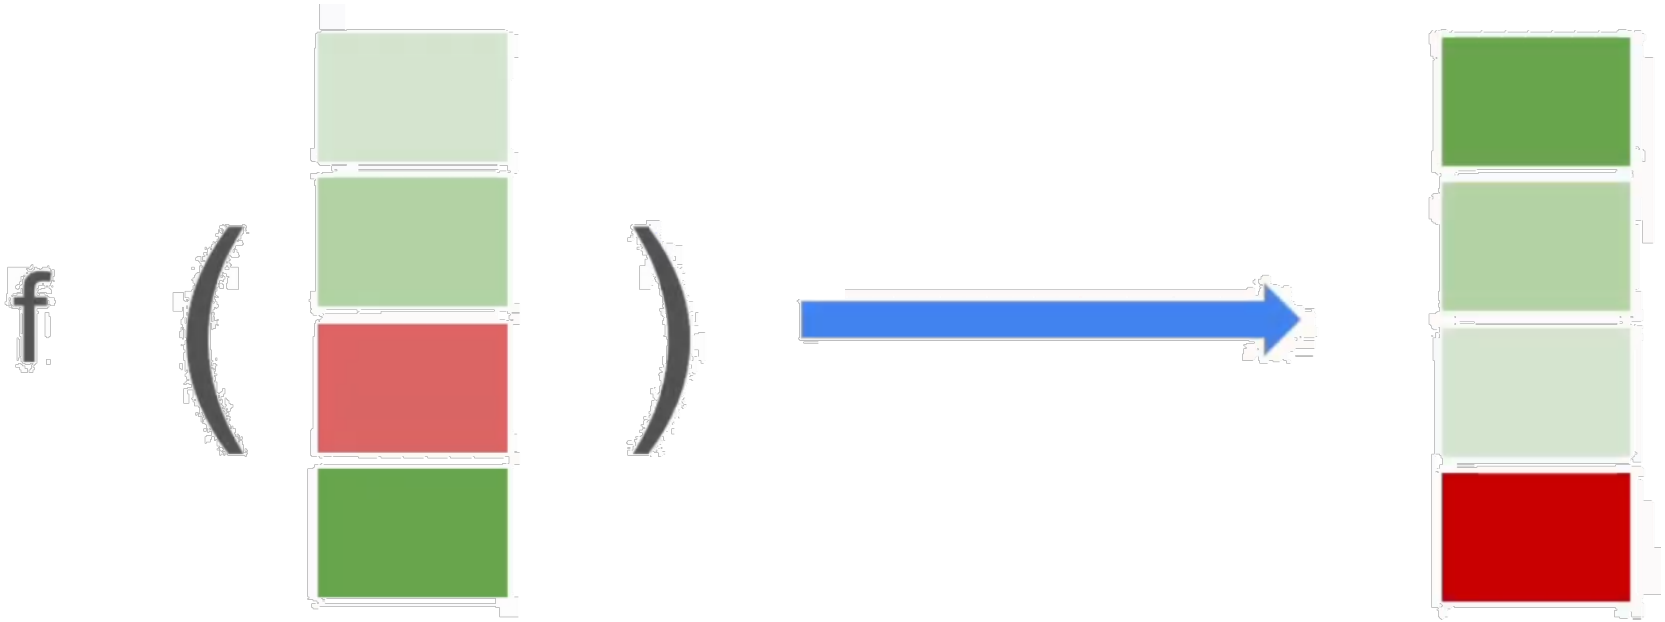

Ranking

amazon.com

Recommended for You

Amazon.com has new recommendations for you based on items you purchased or told us you own.

Google Apps Deciphered: Compute in the Cloud to Streamline Your Desktop

Google Apps Administrator Guide: A Private-Label Web Workspace

Googlepedia: The Ultimate Google Resource (3rd Edition)

Ranking

The Latest React Footwear

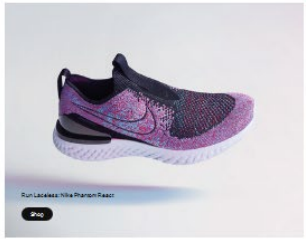

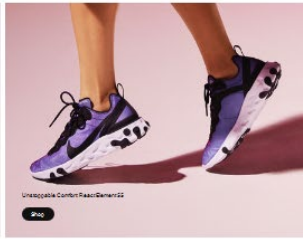

Coming 5.9

Trending

Ranking

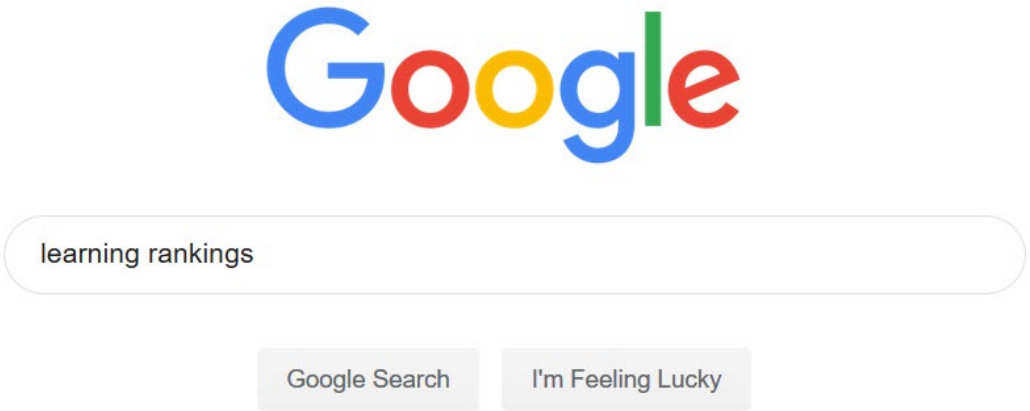

Ranking

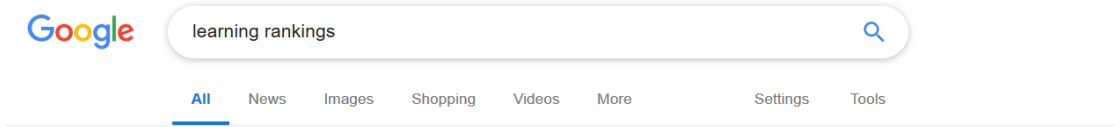
Google
learning rankings
All
News
Images
Shopping
Videos
More
Settings
Tools
About 195,000,000 results (0.39 seconds)
Learning to rank - Wikipedia
https://en.wikipedia.org/wiki/Learning_to_rank
Learning to rank. Learning to rank or machine-learned ranking (MLR) is the application of machine learning, typically supervised, semi-supervised or reinforcement learning, in the construction of ranking models for information retrieval systems.
Discounted cumulative gain · Mean reciprocal rank · Ranking (information retrieval)
You've visited this page 3 times. Last visit: 5/2/19
Best Graduate Education Schools in 2020 - US News
https://www.usnews.com › Education › Grad Schools
See the top ranked education departments and find the best education school for ... reputation and research activity to rank the top graduate education schools.
U.S. News Announces 2019 Best Online Programs Rankings | Press ...
https://www.usnews.com/.../us-news-announces-2019-best-online-programs-rankings
Jan 15, 2019 - ... University is No. 1 for bachelor's, Johns Hopkins is first for nursing and University of Florida ties with Clemson for best in education.

Ranking

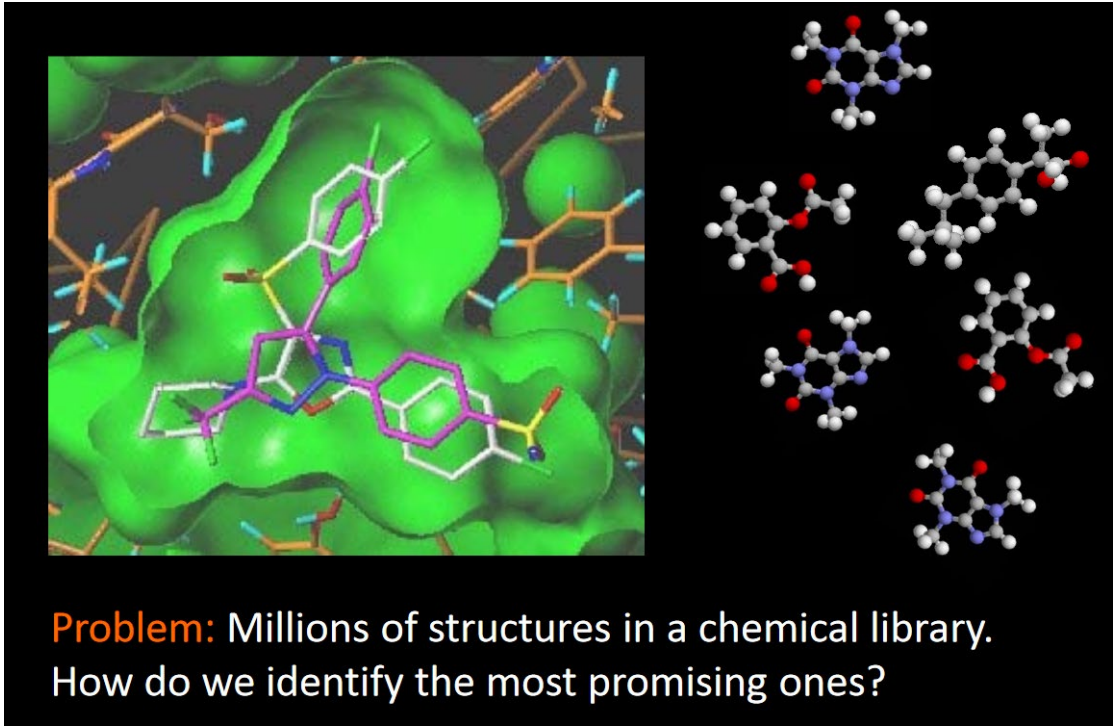
Problem: Millions of structures in a chemical library.
How do we identify the most promising ones?

Ranking for document retrieval

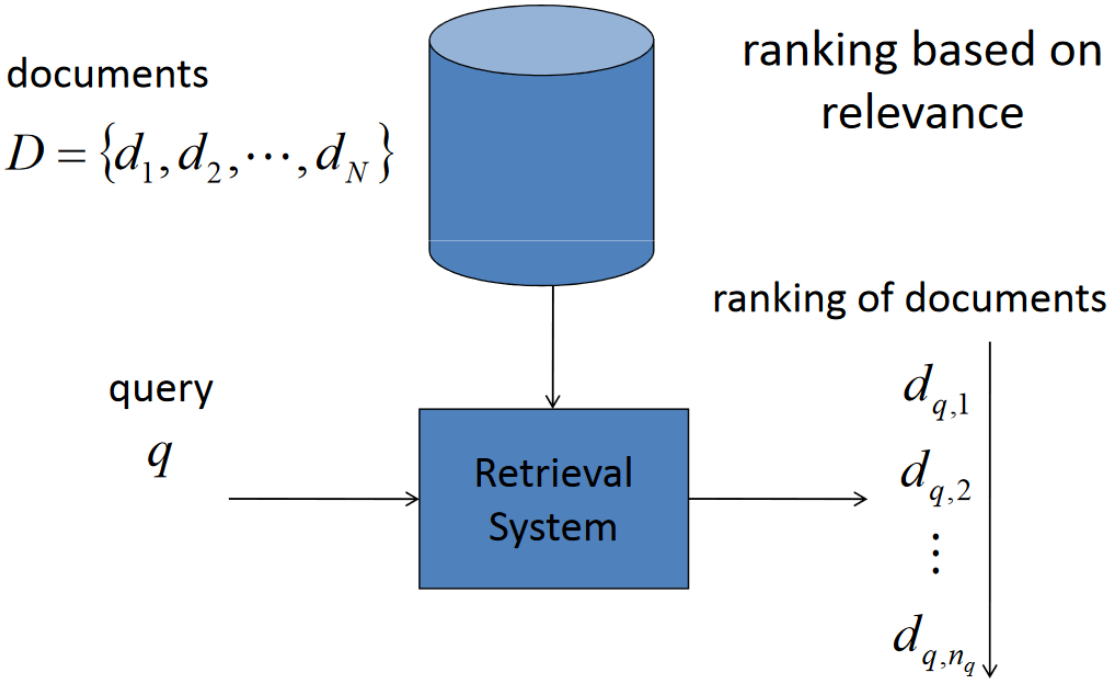

Ranking classifier

- Given
  - Two items d_i and d_j
  - Query q
- Predict
  - Preference
  - $d_i > d_j$ *or* $d_i < d_j$
  - With respect to q

- How do we do this?
  1. Train classifier to predict preferences
  2. Turn predicted preferences into ranking

Simple ranking classifier

Algorithm 17 NAIVERANKTRAIN(*RankingData*, BINARYTRAIN)

```
1:  D ← [ ]
2:  for n = 1 to N do
3:      for all i, j = 1 to M and i ≠ j do
4:          if i is prefered to j on query n then
5:              D ← D ⊕ (x_nij, +1)
6:          else if j is prefered to i on query n then
7:              D ← D ⊕ (x_nij, −1)
8:          end if
9:      end for
10: end for
11: return BINARYTRAIN(D)
```

Simple ranking classifier

Algorithm 18 NAIVERANKTEST(f, $\hat{x}$)

1: $score \leftarrow \langle 0, 0, \ldots, 0 \rangle$ // initialize M-many scores to zero
2: **for all** $i, j = 1$ **to** M **and** $i \neq j$ **do**
3: $y \leftarrow f(\hat{x}_{ij})$ // get predicted ranking of i and j
4: $score_i \leftarrow score_i + y$
5: $score_j \leftarrow score_j - y$
6: **end for**
7: **return** ARGSORT($score$) // return queries sorted by score

BIPARTITE

Ranking formalization

- Ranking is function
  - Maps item d_u to position i in list, $i \in 1,2,..,M$
  - $f : d_u \rightarrow \sigma_u$

- If $\sigma_u < \sigma_v$
  - u is preferred to v (u is higher on ranked list)

Loss function

- $\sum_{u \neq v} [[\sigma_u < \sigma_v] [\hat{\sigma}_v < \hat{\sigma}_u] \omega(\sigma_u, \sigma_v)]$

TASK: ω-RANKING

Given:

1. An input space $\mathcal{X}$
2. An unknown distribution $\mathcal{D}$ over $\mathcal{X} \times \Sigma_M$
3. A training set D sampled from $\mathcal{D}$

Compute: A function $f : \mathcal{X} \to \Sigma_M$ minimizing:

$$\mathbb{E}_{(x,\sigma)\sim\mathcal{D}} \left[\sum_{u \neq v} [\sigma_u < \sigma_v] \, [\hat{\sigma}_v < \hat{\sigma}_u] \, \omega(\sigma_u, \sigma_v) \right] \tag{6.7}$$

where $\hat{\sigma} = f(\boldsymbol{x})$

Cost function

- $\omega(i, j)$
- Kemeny distance measure

Ranking classifier with cost function

Algorithm 19 RANKTRAIN($\mathbf{D}^{rank}$, ω, BINARYTRAIN)

1: $\mathbf{D}^{bin} \leftarrow [\,]$
2: **for all** $(x, \sigma) \in \mathbf{D}^{rank}$ **do**
3: **for all** $u \neq v$ **do**
4: $y \leftarrow \text{SIGN}(\sigma_v - \sigma_u)$ // y is +1 if u is prefered to v
5: $w \leftarrow \omega(\sigma_u, \sigma_v)$ // w is the cost of misclassification
6: $\mathbf{D}^{bin} \leftarrow \mathbf{D}^{bin} \oplus (y, w, x_{uv})$
7: **end for**
8: **end for**
9: **return** BINARYTRAIN($\mathbf{D}^{bin}$)

How can we use instance weights?

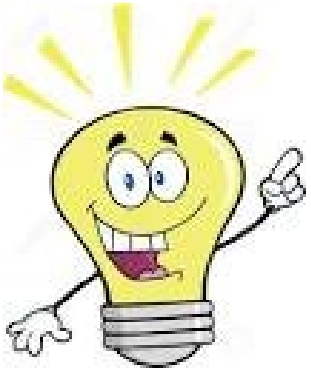

Old (binary tree):

$$Gain(S, A) = Entropy(S) - \left(\left|S_{left}\right| \times Entropy_{left}\right) - \left(\left|S_{right}\right| \times Entropy_{right}\right)$$

New (binary tree):

$$Gain(S, A) = Entropy(S) - \left(\left|S_{left}\right| \times \omega_{left} \times Entropy_{left}\right) - \left(\left|S_{right}\right| \times \omega_{right} \times Entropy_{right}\right)$$

New Ranking will use QuickSort

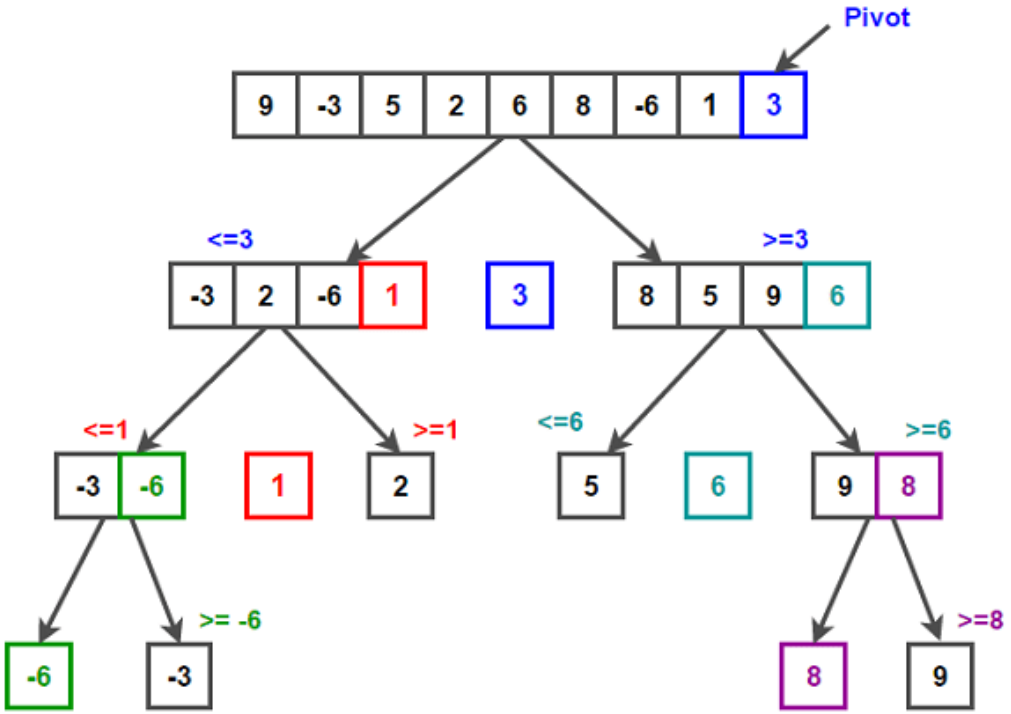

Ranking classifier with cost function

Algorithm 20 RANKTEST(f, $\hat{x}$, obj)

```
1:  if obj contains 0 or 1 elements then
2:      return obj
3:  else
4:      p ← randomly chosen object in obj                    // pick pivot
5:      left ← [ ]                                           // elements that seem smaller than p
6:      right ← [ ]                                          // elements that seem larger than p
7:      for all u ∈ obj \{p} do
8:          ŷ ← f(x_up)                                      // what is the probability that u precedes p
9:          if uniform random variable < ŷ then
10:             left ← left ⊕ u
11:         else
12:             right ← right ⊕ u
13:         end if
14:     end for
15:     left ← RANKTEST(f, x̂, left)                          // sort earlier elements
16:     right ← RANKTEST(f, x̂, right)                        // sort later elements
17:     return left ⊕ ⟨p⟩ ⊕ right
18: end if
```

New Ranking will use Probabilities

Classification

A: 5 Ways to make a million dollars without working

B: 10 Reasons you Should Drink Milk Every Morning (you won't believe number 7!)

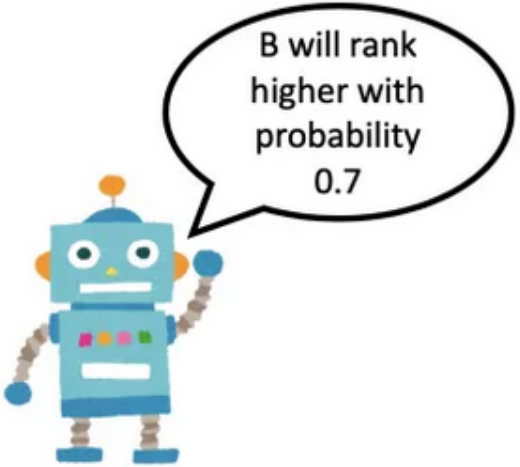

Learning to Rank

Homework Preview

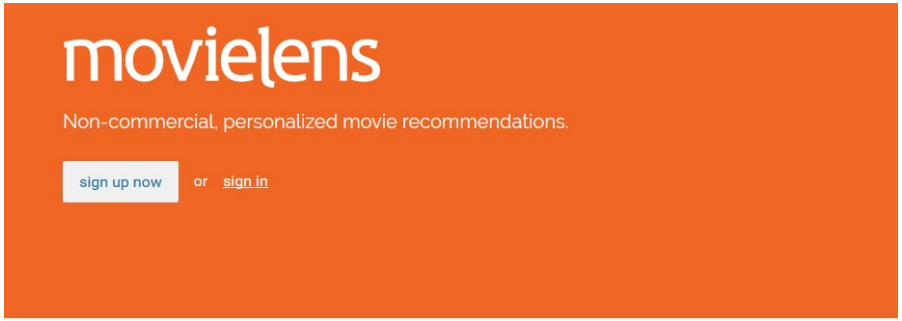

event_1: <customer_1, movie_1, fail>
event_2: <customer_1, movie_2, fail>
event_3: <customer_1, movie_3, success>
event_4: <customer_2, movie_2, fail>
event_5: <customer_2, movie_3, success>
…

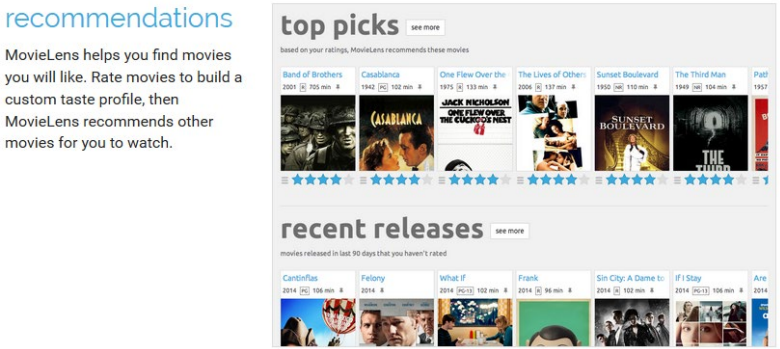

{release_date,Action,Adventure,Children's,Comedy,
Crime,Documentary,Drama,Fantasy,Film-Noir,Horror,
Musical,Mystery,Romance,Sci-Fi,Thriller,War,Western,
ratings_average,ratings_count,price}